TrustLayer

The API apps call to know if the other side is a real human — not a deepfake, bot, or synthetic identity — without forcing users to give up privacy.

What we're solving

The internet was built on assumptions AI has broken: a face, a voice, a conversation, an account — none of these prove humanity anymore. Generative AI has driven the cost of impersonating a human toward zero: deepfakes, voice clones, synthetic reviews, AI-assisted interviews, and automated fraud are now cheap and scalable. Today's trust tools are fragmented point solutions — CAPTCHA stops cheap bots, KYC verifies a document once, fraud tools flag transactions after the fact. None answer the question applications actually need, continuously: *"Can I trust this session, right now?"* TrustLayer exists to build that missing trust infrastructure layer for the AI-native internet.

The product

TrustLayer is two pieces: **TrustLayerOS**, the production control plane that scores whether a digital session is a live human — not a deepfake, replay, empty bot, or synthetic identity — and **TrustLayerSDK**, the client you call before you trust the other side.
The SDK captures session signals (behavior, device, network, camera, microphone under user consent), sends them to TrustLayerOS, and returns a decision your product can enforce: **allow · monitor · verify · review · block**, with *human_probability*, *deepfake_risk*, and *integrity_score*. The first wedge is interview / live-call authenticity — prove the person on the call is the real candidate, live and unassisted — then expand to transactions, fraud, bots, and AI agents.

1 · Capture	2 · Score	3 · Decide
SDK collects behavior, device, network, media signals with consent	TrustLayerOS runs trust + risk engines across interview, deepfake, bot, fraud, anomaly modules	Policy engine returns allow / monitor / verify / review / block with field provenance

Why now

Generative AI and autonomous agents have changed the economics of deception: convincing fraud no longer needs expertise or resources — it scales. Video and voice are approaching human-level realism, so traditional trust signals are unreliable. Just as SSL secured websites, cloud secured infrastructure, and identity providers secured authentication, the next decade needs a layer that secures digital interactions. TrustLayer is that layer — privacy-first (consent, data minimization, features over raw recording), developer-first (simple API, open-source SDKs in TypeScript and Python), and built for continuous trust, not one-time IDV.

Not KYC. Not CAPTCHA. Not proof-of-personhood.

CAPTCHA / Turnstile	Stops cheap bots at the edge — not live video authenticity
Persona / Onfido	Verifies a document once — not continuous during the call
Worldcoin-style PoP	Proves unique personhood — not session risk
TrustLayer	Continuous multi-signal session authenticity: capture → score → decide

Help us get into Y Combinator

We're applying to YC to build the trust layer of the AI internet. If you believe the internet needs a programmable way to know there's a real human on the other side of a call, message, or transaction — please refer us. In an era of AI and deepfakes, protecting people starts with knowing who — and what — is really on the other end of the line.

GitHub: TrustLayerSDK · TrustLayerOS · TrustLayerDocs

One-liner: "Can I trust this interaction?"